

1 Caches

- **Temporal locality:** Use it recently, want to use it again soon
 - Keep most recently accessed data closer
- **Spatial locality:** Use a piece of memory, want to use neighboring pieces soon
 - Move blocks of contiguous words closer to processor

1.1 TIO

- **Tag:** Check if have correct block
 - Num bits = mem addr bits $- I - O$
 - Distinguish diff blocks that use same index
- **Index:** Select block
 - Num bits = $\log_2(\# \text{ of indices})$
- **Offset:** Byte offset within block
 - Num bits = $\log_2(\text{size of block})$
- **Valid bit:** 1 for valid entries
- Check in order of Index, Valid, Tag, Offset
- Cache size = $I * O$
- Address bit width = $\log_2(\text{memory size}) = T + I + O$

1.2 Policies

- **Write-through:** Update cache and memory at same time
- **Write-back:** Update word in cache block, allowing memory word to be stale
 - Add dirty bit to block
 - Update memory when block is removed (and dirty bit = 1)
- **Write-allocate:** Bring block into cache after write miss

1.3 Cache Types

- **Direct-mapped cache:** (1 way set assoc.)
 - **Replacement:** Index specifies position to replace
 - Each memory addr associated with 1 block in cache
 - Look in single location in cache for data
- **Fully Associative Caches:** (M way set assoc.)
 - **Replacement:** Can replace any position
 - Tag and offset same as before, no index now
 - **Benefits:** No conflict misses since data can go anywhere
 - **Drawbacks:** Hardware is difficult to make
- **Set Associative Cache:**
 - **Replacement:** Can replace any within the set
 - Tag and offset same as before, index points to correct set
 - Each set contains multiple blocks
 - Once found correct set, compare with all blocks in set
 - * Compare tag with all tags in set
 - **Benefits:** Avoids conflict misses using N comparators

1.4 Cache Misses

- 1) **Compulsory:**
 - First time you access a certain block
 - **Solution:** bigger block size
 - 2) **Conflict:**
 - Wouldn't have occurred with fully associative cache (LRU)
 - **Solution:** increase associativity or replacement policy
 - 3) **Capacity:**
 - Independent of associativity of cache
 - If not compulsory or conflict miss
 - Occurs you still miss with fully associative cache (LRU)
 - **Solution:** Increase cache capacity
- **Replacement policies:**
 - **Least Recently Used (LRU):**
 - * Better temporal locality, harder to keep track of
 - **First In First Out (FIFO):** Just track initial order
 - **Random:** Good for data with low temporal locality

1.5 AMAT

- **Average Memory Access Time (AMAT):**
 - $\text{AMAT} = \text{Hit Time} + \text{Miss Rate} * \text{Miss Penalty}$
- **Global miss rate:** misses / total # accesses for cache system
- **Local miss rate:** misses / total # accesses for cache level
- Recursive expression for multi-level cache

2 Virtual Memory

- **Virtual address:** Can be diff # bits than physical
- **Physical address:** Can be diff # bits than physical
- **Page Table:** Stores mapping of VPN to PPN
 - Also includes valid bit, dirty bit, permission bits
 - Page Table entry: a row of the PT
 - Page Table Base Register (PTBR) points to PT
- Hierarchical PT has many page numbers for L1, L2 caches
- **Protection Fault:** PT Entry permission bits prohibit op
- **Page fault:** PT Entry has valid bit set to false (not in mem)
 - Transfer page to memory (evict a page if necessary)
 - Add mapping to PT and TLB

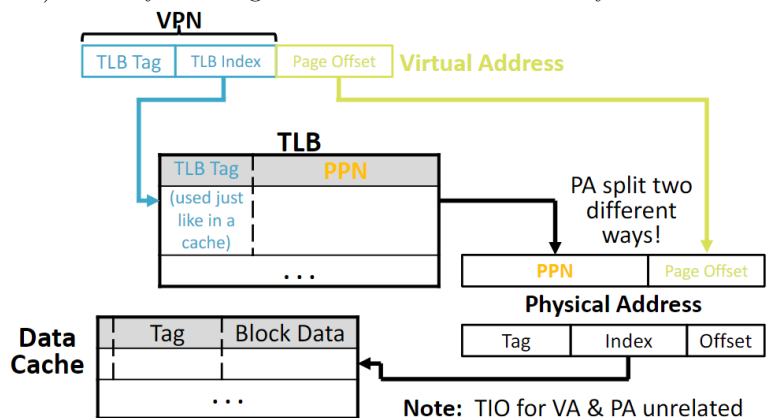
2.1 TLB Address Lookup

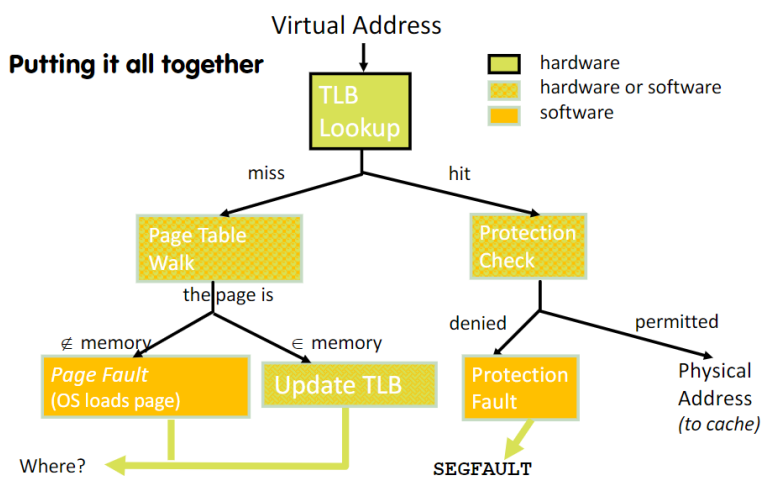
- **Translation Lookaside Buffer (TLB):** Cache for PT
 - Each block is a single PT entry
 - **TLB Miss:** If entry is not in TLB
 - TLB is a subset of the PT
 - **TLB Reach:** Size of largest virtual address space that can fit in TLB

TLB Valid	TAG (VPN)	Page Table Entry		
		Page Dirty	Perm Bits	PPN
– TLB Entry –				
– TLB Entry –				

– Procedure:

- 1) Get VPN from Virtual Address
 - Translate VPN to PPN using TLB
- 2) If TLB does not have VPN, check if PT contains
- 3) If PT does not have VPN, this is a page fault
 - Fetch physical page from disk to memory
- 4) Update PT Entry, load entry into TLB
- 5) Use Physical Page and offset to access memory





3 IO

- **Polling:** Forces hardware to wait on ready bit
 - 1) Processor reads from control register in loop until ready
 - 2) Processor does IO, resets control register bit to 0
 - **Pros:** Low latency, low overhead when data is available
 - * Good for when you need reply to make progress
 - **Cons:** Can't do anything else (or sleep) while polling
- **Interrupts:** Hardware fires exception when ready
 - CPU changes PC to execute interrupt handler
 - Low data rate (mouse, keyboard): use interrupts
 - High data rate (network, disk): start with interrupts, switch to DMA
 - **Pros:** Can do work while waiting, wait for many at once
 - * Good for devices that take long to respond
 - **Cons:** Nondeterministic when interrupt occurs, higher latency, worse throughput

3.0.1 DMA

- **Direct Memory Access (DMA):** Allows IO devices to directly read/write main memory
- **Incoming Data:**
 - 1) Receive interrupt from device
 - 2) CPU takes interrupt, initiates transfer
 - Instructs DMA to place data at certain address
 - 3) DMA handles transfer as CPU does other things
 - 4) DMA interrupts CPU once finished
- **Outgoing Data:**
 - 1) CPU initiates transfer, confirms device is ready
 - 2) CPU begins transfer
 - Instructs DMA that data is at certain address
 - 3) DMA handles transfer as CPU does other things
 - 4) DMA interrupts CPU once finished

3.1 OS

- When to transition into supervisor mode:
 - Interrupts: Something external to the running program
 - * Async to current program (ex. key press, disk IO)
 - Exception: Caused by an event during execution of an instruction by the running program
 - * Sync, on instruction that causes exception (ex. mem error, bus error, illegal instruction)
 - Trap: action of servicing interrupt or exception by hardware jump to interrupt/trap handler code
- Ecall: Trigger exception to higher privilege
- EbreaK: Trigger exception within current privilege

4 Data Parallelism

- **Instruction level parallelism:** Pipelining
- **Data level parallelism:** Single instruction operation on multiple streams of data
- **Request level parallelism:** Many requests/sec reading a database (not read/write sync)
- Flynn's Taxonomy:
 - SISR: Single Instruction Single Data (traditional programs)
 - SIMD: Single Instruction, Multiple Data (SSE intrinsics)
 - MISD: Multiple Instruction, Single Data
 - MIMD: Multiple Instruction, Multiple Data (map reduce)

4.1 SIMD

- Example methods:

```

1 _mm_loadu_si128((__m128i*) (vals + offset));
2 _mm_storeu_si128((__m128i*) p, _sum_vector);

```

- Tail Case:

```

1 for(int i = N / 4 * 4; i < N; i++) {
2     // Tail Case
3 }

```

- Unroll: 4 iterations per time

5 Thread-Level Parallelism

- Often more logical CPUs than physical CPUs
- OpenMP uses a fork-join model

- **Amdahl's Law:** True Speedup = $\frac{1}{S + \frac{1-S}{P}}$
 - S is portion that is serial
 - P is the speedup factor
 - $\lim_{P \rightarrow \infty} \frac{1}{S + \frac{1-S}{P}} = \frac{1}{S}$

5.1 OpenMP Directives

- Import: `#include <omp.h>`
- Parallel: `#pragma omp parallel`
 - Each thread runs a copy of the code in block
- For: `#pragma omp parallel for`
 - Each thread runs different iterations of the for loop
- Thread functions:
 - `int omp_get_thread_num()`: thread num executing the code
 - `int omp_get_num_threads()`: total num of threads
- Critical: `#pragma omp critical`
 - Only 1 thread can run critical part at a time
- Reduction: `#pragma omp parallel for reduction(op:var)`
 - Reduces thread results into var using operation

5.2 Synchronization

- Use locks to control access to shared resources
- **Critical section:** Area that must be serialized
- **Solution:** Atomic read/write in single instruction
 - **Atomic Memory Operation (AMO):** Atomic swap of register and memory
 - * **Load-reserve, store-conditional:** Uninterrupted execution across multiple inst
 - * **Amo.swap:** Uninterrupted memory ops in 1 instruction
- **OpenMP locks:**
 - Declare lock: `omp_lock_t lock`
 - `omp_init_lock(&lock)`, `omp_set_lock(&lock)`
 - `omp_unset_lock(&lock)`, `omp_destroy_lock(&lock)`

Compare-and-swap

```

Start: addi      t0 x0 1 # Locked = 1           # a0 holds address of memory location
      amoswap.w.aq t1 t0 (a0)                   # a1 holds expected value
      bne        t1 x0 Start                   # a2 holds desired value
# If the lock is not free, retry                # a0 holds return value, 0 if successful, !=0
                                              # otherwise

      ... # Critical section

      amoswap.w.rl x0 x0 (a0) # Release lock

cas:
      lr.w t0, (a0) # Load original value.
      bne t0, a1, fail # Doesnt match, so fail.
      sc.w a0, a2, (a0) # Try to update.
      jr ra # Return.

fail:
      li a0, 1 # Set return to failure.
      jr ra # Return.

```

5.3 Cache Coherency

- Cache-incoherence: Each hardware thread cache does not have up to date memory
- Read-modify-write: Interleaving of R M W stages
- Cache states: MOESI
 - 1) **Modified:** Up to date data, changed/dirty
 - o No other cache has a copy, ok to write
 - o Memory out of date (need to write back)
 - 2) **Owner:** Up to date data
 - o Other caches can have copy, but must be shared
 - o Variation of Shared, supplies data instead of going to memory on read miss
 - o Can become Shared after writing modifications to memory
 - o Variation of Modified, avoids writing to memory on miss
 - o Can become Modified if all shared are invalidated
 - 3) **Exclusive:** Up to date data, unchanged
 - o No other cache has a copy, ok to write
 - o Memory up to date (don't need to write back)
 - 4) **Shared:** Up to date data
 - o Other caches can have a copy
 - 5) **Invalid:** Not in cache
- Blocks move to Invalid if another processor writes to it
- Blocks move to Shared if another processor reads it
- **False sharing:** Block goes between 2 caches even though variables are disjoint
 - o **Coherence/communication miss:** Disjoint variables, but have to invalidate block as one updates
- Permitted cache states:

	M	O	E	S	I
M					Y
O				Y	Y
E					Y
S		Y			Y
I	Y	Y	Y	Y	Y

- **MOESI FSM:**
 - o Start at invalid
 - o Transition based on the processor reading/writing
 - o **Read/Write Hits:** The processor that is caching the block
 - o **Probe Read/Write Hit:** A different processor attempts to access a block that is cached by another processor.

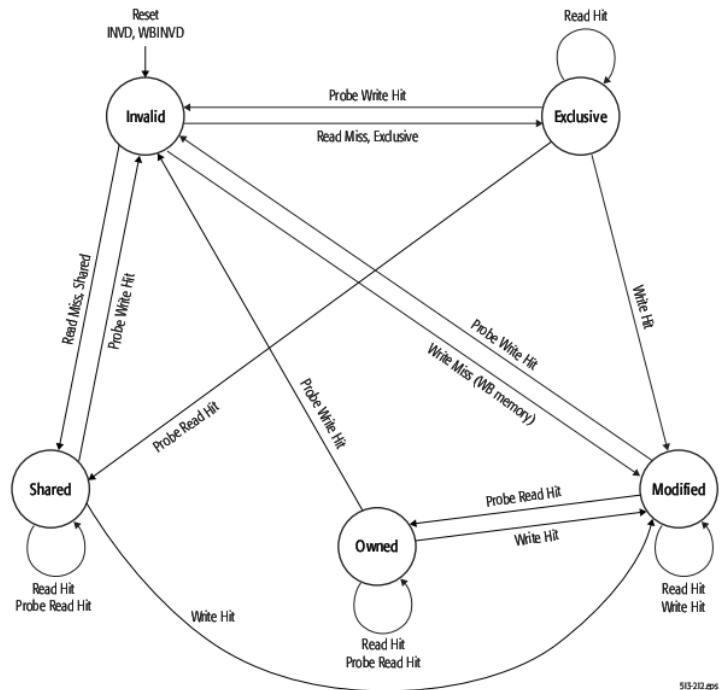
6 Distributed Computing

6.1 MapReduce

- **Map:** processes input key/value pair
 - o Slices data into shards/splits, distributed to workers
 - o Produces set of intermediate pairs
 - * Use `emit(key k, value v)` to output K-V pair (k, v)
- Group by key, send to reduce
- **Reduce:** Combines intermediate values for a particular key, produces an output value
- All Map must finish before reduce, but can start reading values

6.2 Spark

- Much faster than MapReduce
- Resilient Distributed Dataset (RDD): Abstraction of a DB
- Transforms (RDD \rightarrow RDD):
 - o **map(f)** : Calls f on each source element
 - o **flatMap(f)** : Each input item can be mapped to ≥ 0 output items
 - o **reduceByKey(f)** : Aggregate by key using $f: (V, V) \rightarrow V$
- Actions (RDD \rightarrow RDD):
 - o **reduce(f)** : Aggregate elements regardless of key
- Example:
 - o `coinData = sc.parallelize(coinPairs)`
 - o `out = coinData.map(lambda (k1, k2): ((k1, k2), 1))`
 - * `.reduceByKey(lambda v1, v2: v1 + v2)`
 - * Computes # coins of each type of each person
- Lazy evaluation model, evaluates after `.collect()`



7 Dependability

- Power Usage Effectiveness (PUE): Total building power / IT equipment power
 - 1.0 is perfect
- Rack contains 40-80 servers, Array contains 16-32 racks

7.0.1 Dependability Measures

- **Mean Time To Failure (MTTF)**: Reliability
- **Mean Time To Repair (MTTR)**: Service interruption
- **Mean Time Between Failures (MTBF)**: $MTBF = MTTF + MTTR$
- **Availability**: $MTTF / (MTTF + MTTR)$
 - Number of 9s availability per year: 1 nine: 90%, 2 nine: 99%

7.0.2 Reliability Measures

- **Annualized Failure Rate (AFR)**: Average number of failures per year
 - $AFR = 24 * 365 / MTTF$
- **Failures in Time (FIT)**: Num of failures expected in 1 billion device-hours of operation
 - $MTBF = 1 \text{ billion} / FIT$

7.1 RAID

- Always know which disk fails
- RAID0 provides maximum usable disk space
- RAID1 is most expensive but highest availability
 - RAID1 writes are faster than RAID5 since no parity

	Config	Pros	Cons
0	Split data across multiple disks	No overhead, fast read/write	Reliability
1	Mirrored disks: extra copy of data	Fast read/write, fast recovery	High overhead = expensive
2	Bit-level striping, 1 disk per parity group	Smaller overhead	Redundant check disks
3	Byte-level striping, single parity disk	Smallest overhead to check parity	Need to read all disks to detect errors
4	Block-level striping, single parity disk	Higher throughput for small reads	Slow small writes (single check disk)
5	Block-level striping, parity distributed	Higher throughput for small writes	Time to repair disk is long

7.2 ECC

- Error Correction Codes (ECC): Protect soft errors
- **Hamming distance**: Diff in # of bits for valid codewords
- **Parity bit**: Forces stored words to have even parity
 - Placed at binary positions of powers of 2
 - XOR all 1 means error
 - Cannot detect an even number of errors
- **Detection**: bit pattern fails codeword check
- **Correction**: map to nearest valid code word
- Hamming ECC uses 1 index left to right (MSB)

Bit position	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
Encoded data bits	p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11	p16	d12	d13	d14	d15
Parity bit coverage	p1	x		x		x		x		x		x		x		x		x		x
	p2		x	x			x	x			x	x		x	x			x	x	
	p4				x	x	x					x	x	x	x					x
	p8								x	x	x	x	x	x	x					
	p16															x	x	x	x	x