

# Pipelining RISC-V

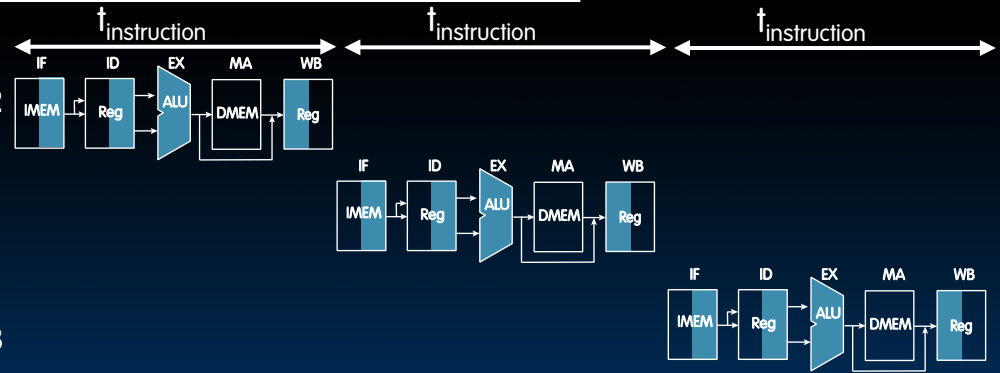
# 'Sequential' RISC-V Datapath

Phase	Pictogram	$t_{step}$ Serial
Instruction Fetch		200 ps
Reg Read		100 ps
ALU		200 ps
Memory		200 ps
Register Write		100 ps
$t_{instruction}$		800 ps

instruction sequence ↓


```

add t0, t1, t2
or t3, t4, t5
sll t6, t0, t3
    
```



# Pipelined RISC-V Datapath

Phase	Pictogram	$t_{step}$ Serial	$t_{cycle}$ Pipelined
Instruction Fetch		200 ps	200 ps
Reg Read		100 ps	200 ps
ALU		200 ps	200 ps
Memory		200 ps	200 ps
Register Write		100 ps	200 ps
$t_{instruction}$		800 ps	1000 ps



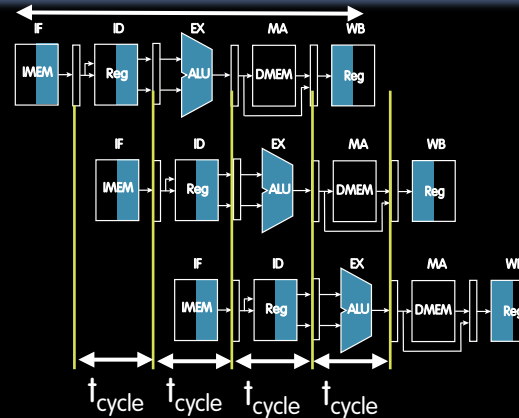
# Pipelined RISC-V Datapath

instruction sequence

add t0, t1, t2

or t3, t4, t5

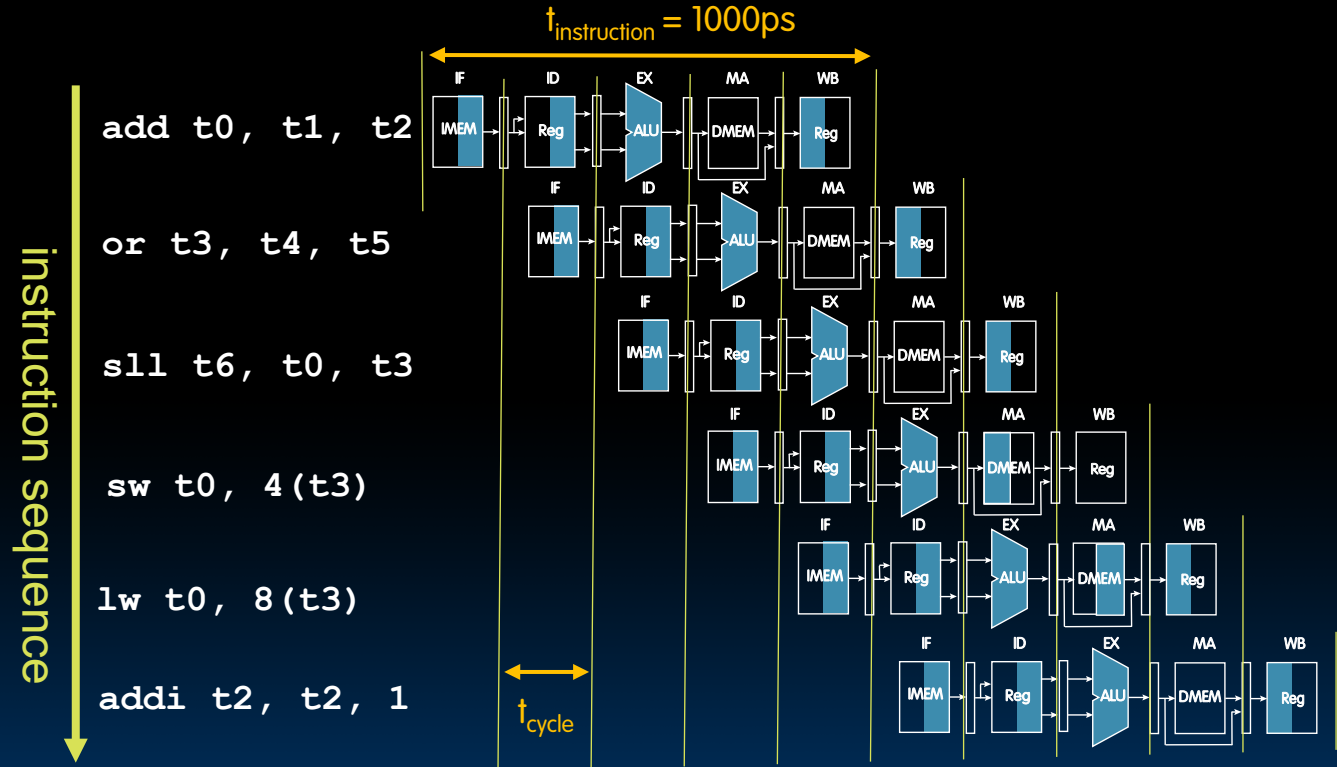
sll t6, t0, t3



	Single Cycle	Pipelined
Timing	$t_{step} = 100 \dots 200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
CPI (Cycles Per Instruction)	$\sim 1$ (ideal)	$\sim 1$ (ideal), $< 1$ (actual)
Clock rate, $f_s$	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = 5 \text{ GHz}$
Relative speed	1 x	4 x

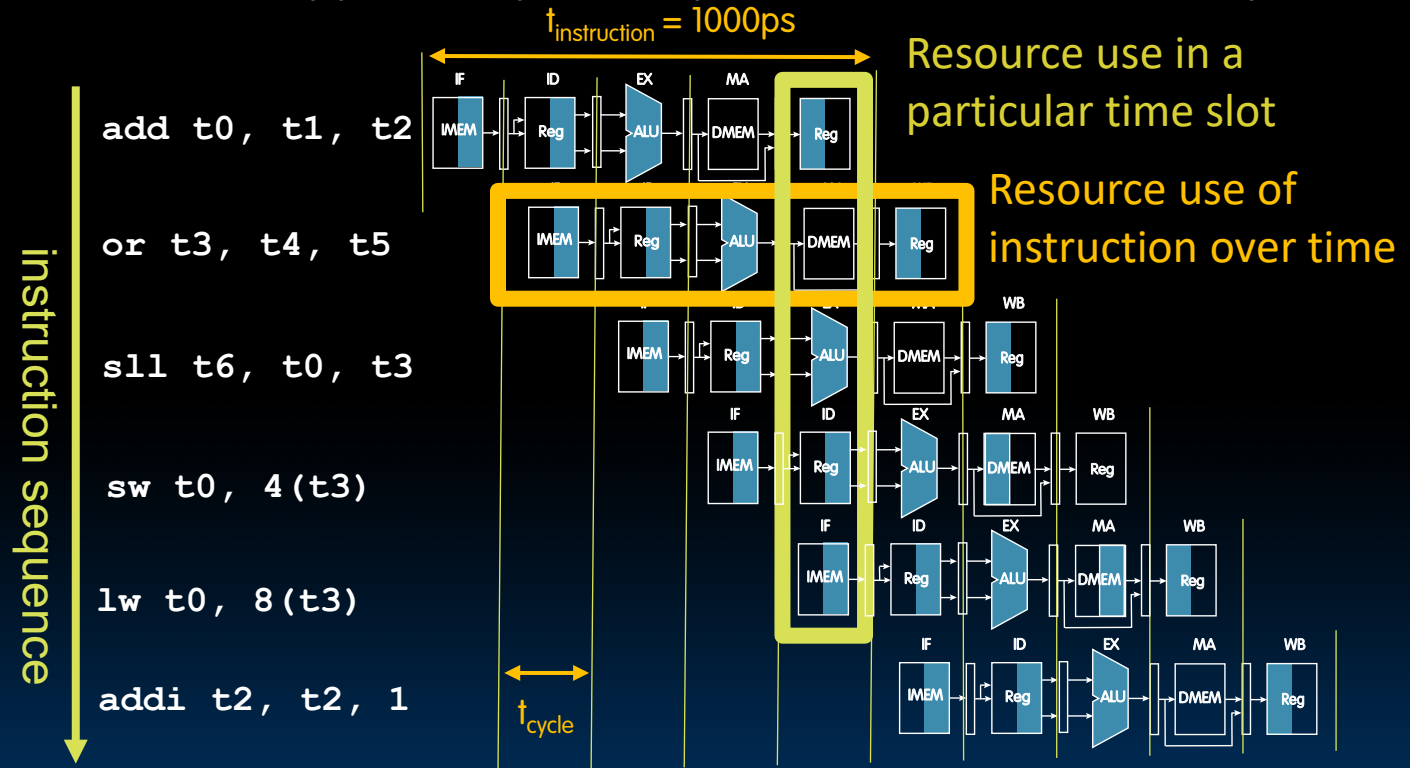
# Sequential vs. Simultaneous

- What happens sequentially and what simultaneously?



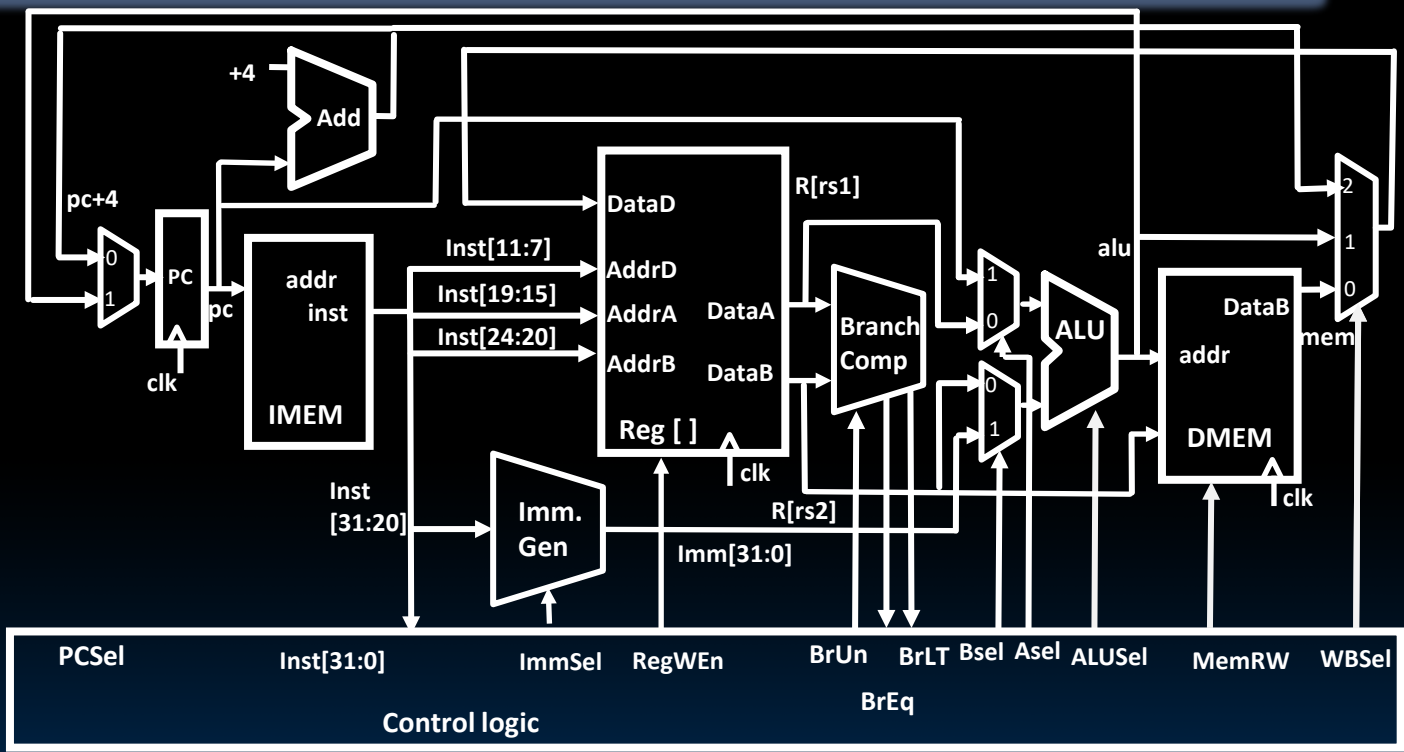
# Sequential vs. Simultaneous

- What happens sequentially and what simultaneously?



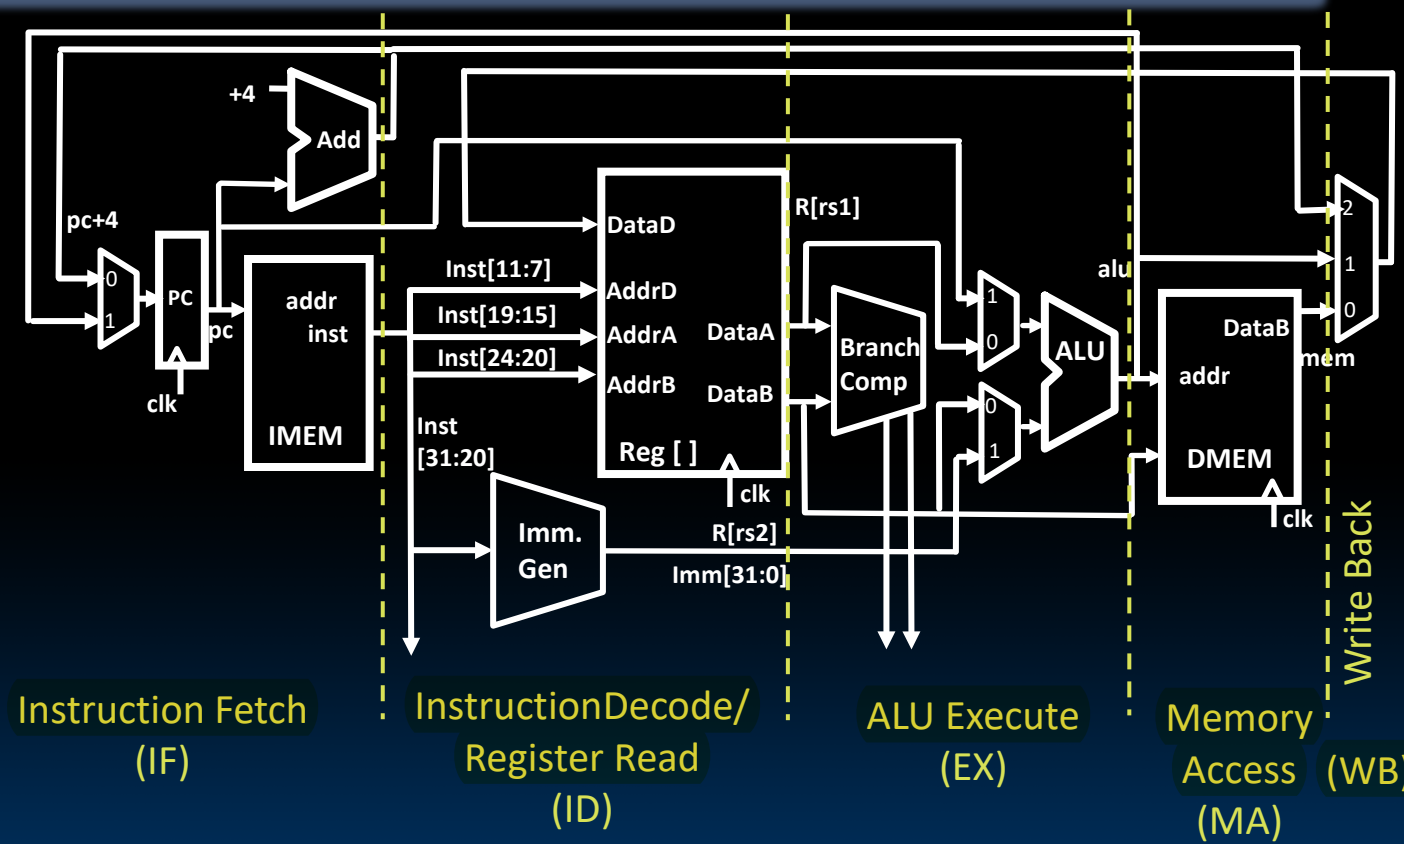
# Pipelining Datapath

# Single-Cycle RV32I Datapath



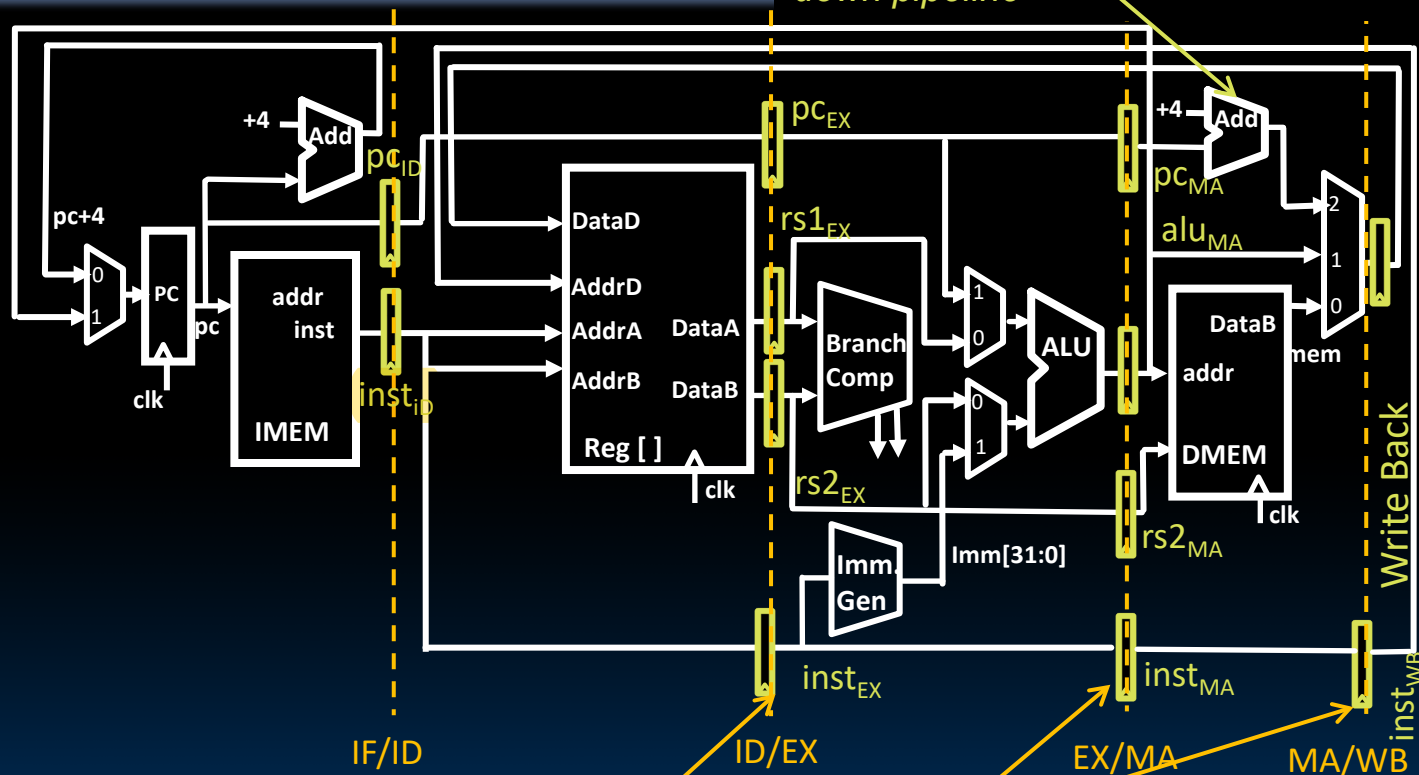


# Single-Cycle RV32I Datapath



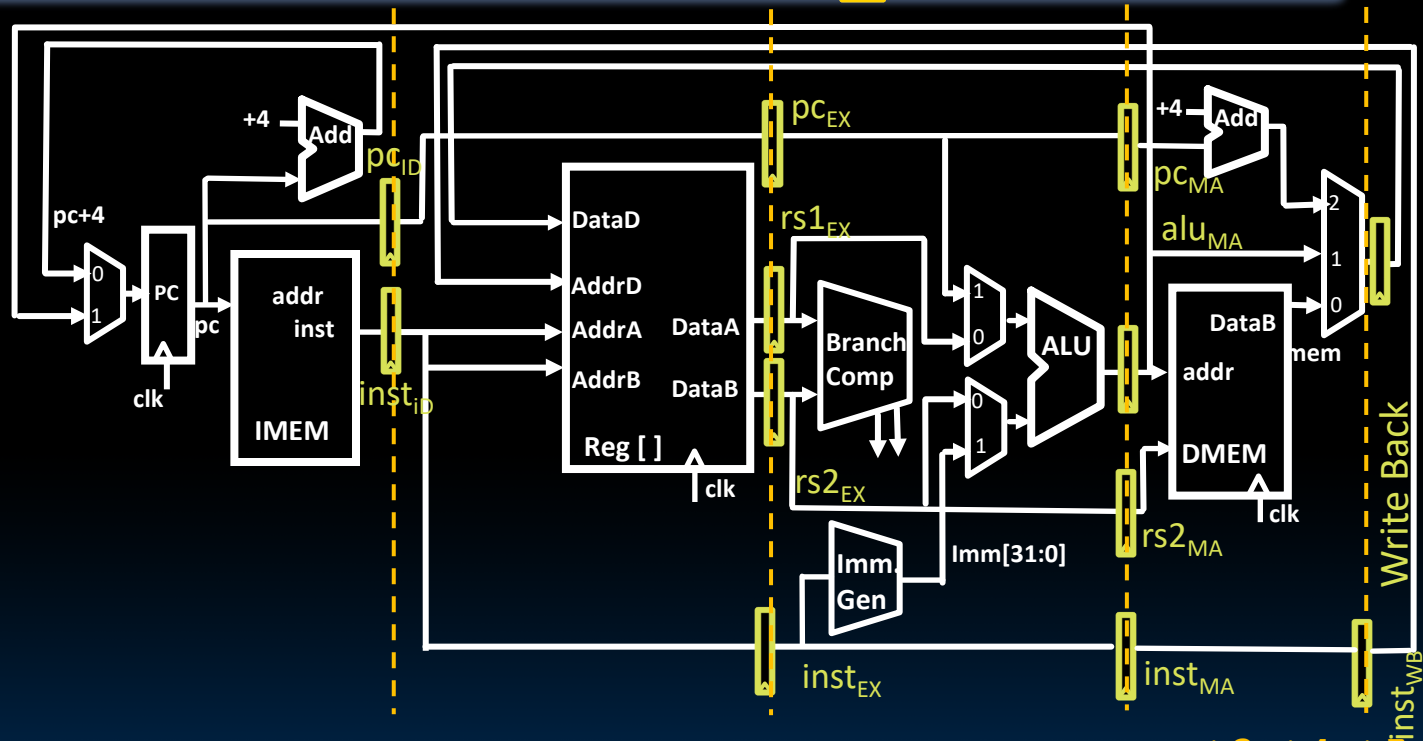
# Pipelined RV32I Datapath

*Recalculate PC+4 in M stage to avoid sending both PC and PC+4 down pipeline*



*Must pipeline instruction along with data, so control operates correctly in each stage*

# Pipelined RV32I Datapath




`lw t0, 8(t3)      sw t0, 4(t3)      slt t6, t0, t3      or t3, t4, t5`

Pipeline registers separate stages, hold data for each instruction in flight

# Pipelined Control

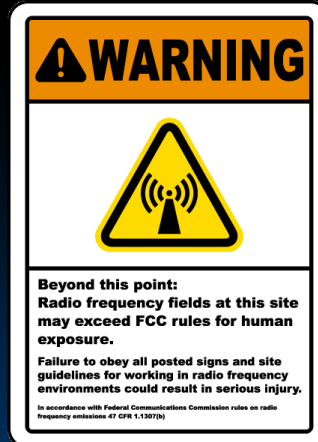


- Control signals derived from instruction
  - As in single-cycle implementation
  - Information is stored in pipeline registers for use by later stages



# Pipeline Hazards

# Hazards Ahead!



# Pipelining Hazards

A *hazard* is a situation that prevents starting the next instruction in the next clock cycle

## 1) *Structural hazard*

- A required resource is busy (e.g. needed in multiple stages)

## 2) *Data hazard*

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

## 3) *Control hazard*

- Flow of execution depends on previous instruction

# Structural Hazard

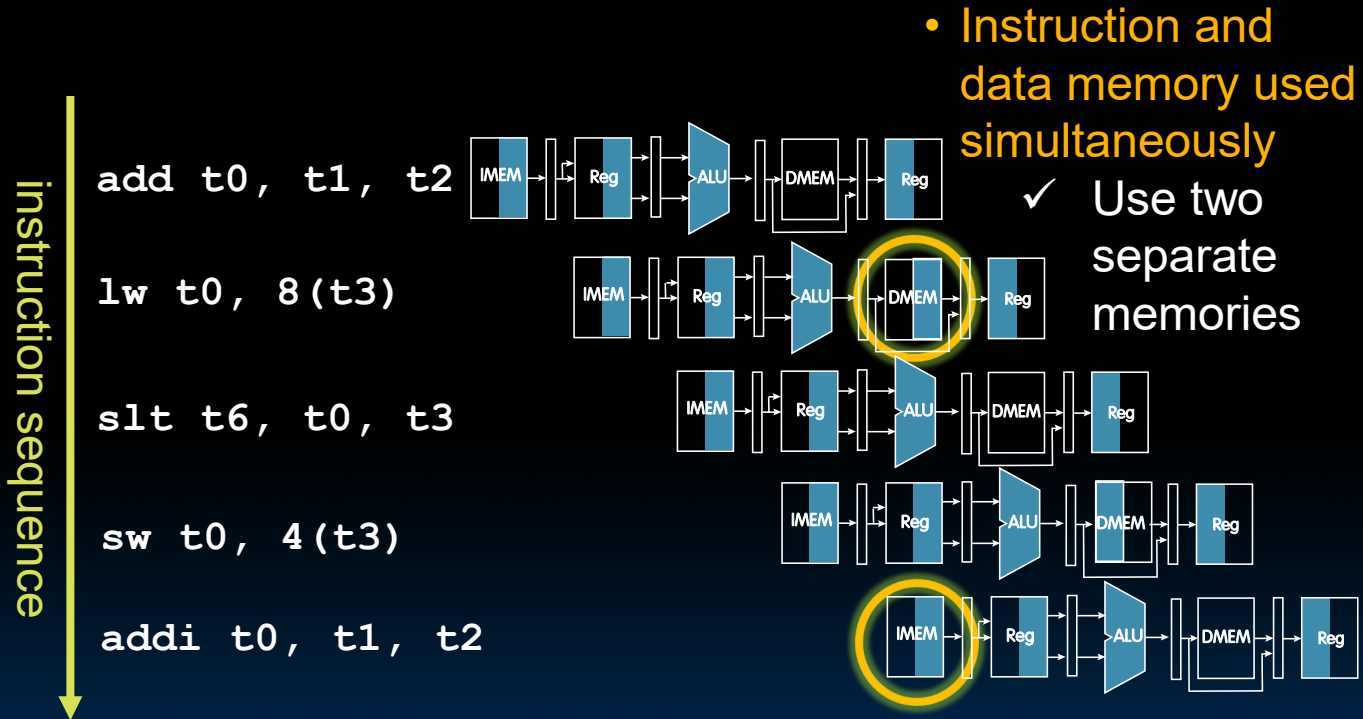
- **Problem:** Two or more instructions in the pipeline compete for access to a single physical resource
- **Solution 1:** Instructions take it in turns to use resource, some instructions have to stall
- **Solution 2:** Add more hardware to machine
- Can always solve a structural hazard by adding more hardware



# Regfile Structural Hazards

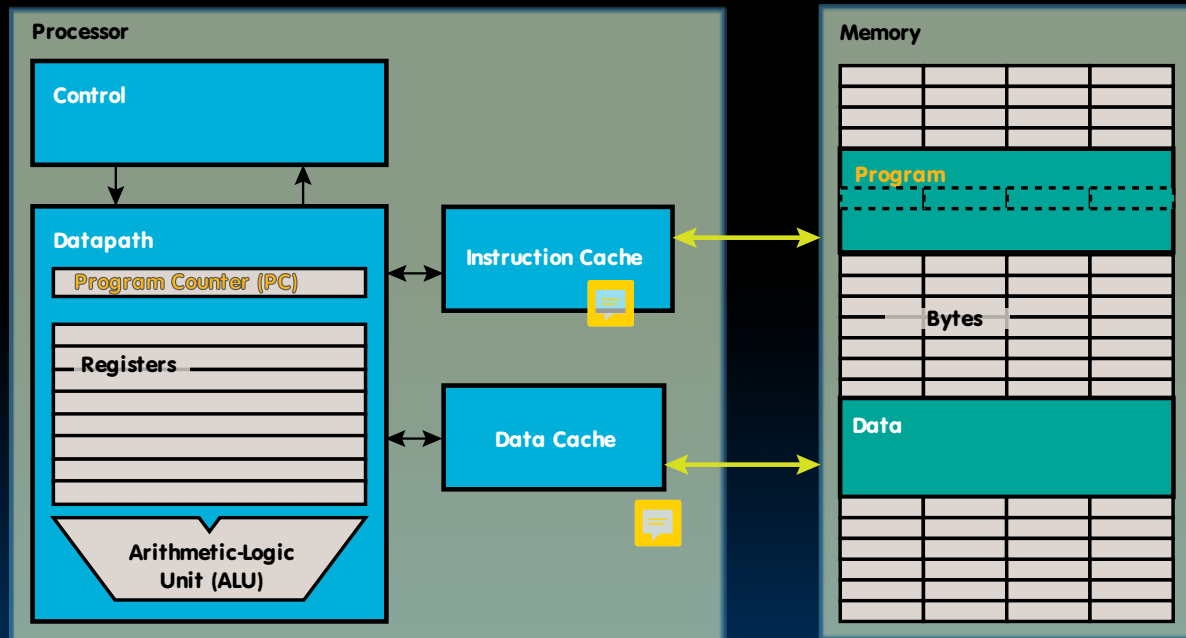
- Each instruction:
  - Can read up to two operands in decode stage
  - Can write one value in writeback stage
- Avoid structural hazard by having separate “ports”
  - Two independent read ports and one independent write port
- Three accesses per cycle can happen simultaneously

# Structural Hazard: Memory Access



# Instruction and Data Caches

- Fast, on-chip memory, separate for instructions and data



# Structural Hazards – Summary

- Conflict for use of a resource
- In RISC-V pipeline with a single memory
  - Load/store requires data access
  - Without separate memories, instruction fetch would have to **stall** for that cycle
    - All other operations in pipeline would have to wait
- Pipelined datapaths require separate instruction/data memories
  - Or separate instruction/data caches
- RISC ISAs (including RISC-V) designed to avoid structural hazards
  - e.g. at most one memory access/instruction



# Data Hazards

# Data Hazard: Register Access


- Separate ports, but what if write to same register as read?

Does **sw** in the




example fetch the

old or new value?




# Data Hazard: Register Access

- Exploit high speed of register file (100 ps)
  - 1) WB updates value
  - 2) ID reads new value
- Indicated in diagram by shading



*Might not always be possible to write then read in same cycle, especially in high-frequency designs. Check assumptions in any question.*

# Data Hazard: ALU Result






# Solution 1: Stalling

- Problem: Instruction depends on result from previous instruction

```
add s0, t0, t1
sub t2, s0, t3
```

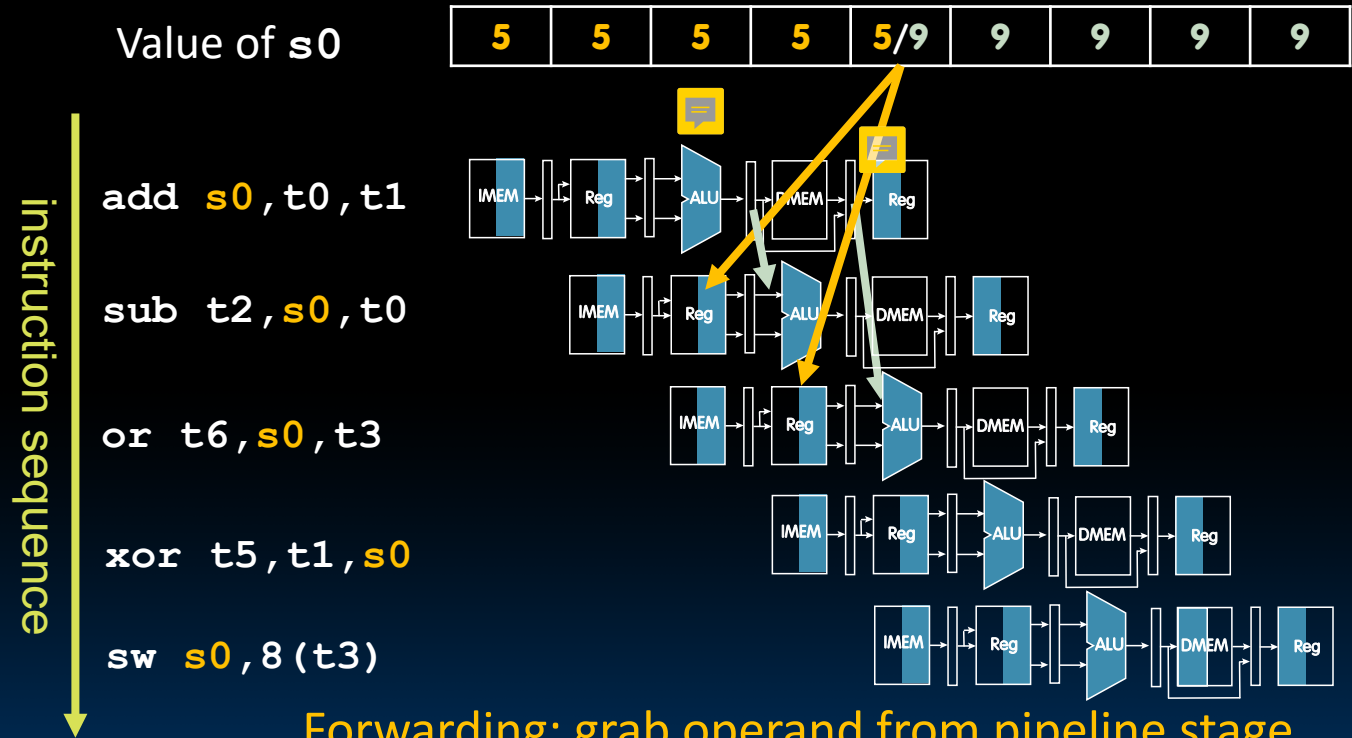


- Bubble:
  - Effectively `nop`: Affected pipeline stages do "nothing"

# Stalls and Performance

- **Stalls** reduce performance
  - But stalls are required to get correct results
- Compiler can arrange code or insert **nops** (`addi x0, x0, 0`) to avoid hazards and stalls
  - Requires knowledge of the pipeline structure

# Solution 2: Forwarding




Forwarding: grab operand from pipeline stage,  
rather than register file

# Forwarding (aka Bypassing)

- Use result when it is computed
  - Don't wait for it to be stored in a register
  - Requires extra connections in the datapath

add s0, t0, t1

sub t2, s0, t3



# Data Needed for Forwarding (Example)

- Compare destination of older instructions in pipeline with sources of new instruction in decode stage.
- Must ignore writes to x0!

add t0, t0, t1

sub t3, t0, t5

sub t6, t0, t3

